

SPP A2 | Floyd Warshal Optimization

- Ayush Sharma
- 2019101004

Floyd-Warshall algorithm is an algorithm for finding shortest paths in a directed weighted graph with positive or negative edge weights. A single execution of the algorithm will find the lengths of shortest paths between all pairs of vertices.

Floyd Warshall Algorithm

Floyd-Warshall algorithm is an algorithm for finding shortest paths in a directed weighted graph with positive or negative edge weights. A single execution of the algorithm will find the lengths of shortest paths between all pairs of vertices.

Optimizations

We were given 10 test cases out of which the largest was having 2229 number of nodes and 155317 number of edges. There can be self loop and multiple edges between same two nodes. Although I profiled all of them but will list timings and cache profile for that largest test case only.

Time taken with the trivial algorithm i.e. no optimisation : ~63 seconds

My machine has total 8 cores

Previous Version (i.e. in Assignment1)

Final Total time:

- On Abacus : ~19.832843 (core independent as no multi-programming was there)
- ~15s on my machine

Process for previous version optimisation

- Converted 2-D array for storing adjacency matrix into 1-D. Then for accessing `A[i][j]` we used pointer notation `(*(A + i*V + j))`, with keeping in mind the fact that generally 2-D matrices are stored in Row Major format in main memory.
- Pointer accessing to memory Restricted pointer access instead of array look-ups.
- Pre increment over post increment Pre-increment is faster than post-increment because post increment keeps a copy of previous (existing) value and adds 1 in the existing value while pre-increment is simply adds 1 without keeping the existing value.
- Used `register` keyword to reduce fetch time for `i`, `j`, `k` iterators which were being used more often.
- Used a temporary variable for storing and reducing memory lookup using pointers i.e.
`register int * kmj = (matrix + k*V + j);` and `register int * imk = (matrix + i*V + k);`
- Tweaked the value of `INFINITY` to 1000000000.
- Removed spurious `if` conditions.

- Used more general memory lookup pointer variable i.e. replace `imk` & `kmj` with `im` and `km`. These new variable will be have more reads / writes ratio as compared to `imk` & `kmj`.

```
void floyd_v2(int * matrix, int V)
{
    register int i,j,k,kmj,imk, v;
    register int * km;
    register int * im;
    v = V;
    for(k=0;k<v;++k)
    {
        km = (matrix + k*v);
        for(i=0;i<v;++i)
        {
            im = (matrix + i*v);
            for(j=0;j+15<v;j+=16)
            {
                imk = (*(im + k));

                kmj = (*(km + j+0));
                if( ((im + j+0)) > ( kmj + imk ) ) ((im + j+0)) = ( kmj + imk );
                ...

                kmj = (*(km + j+15));
                if( ((im + j+15)) > ( kmj + imk ) ) ((im + j+15)) = ( kmj + imk );
            }
            while(j<v)
            {
                kmj = (*(km + j));
                imk = (*(im + k));
                if( ((im + j)) > ( kmj + imk ) ) ((im + j)) = ( kmj + imk );
                j++;
            }
        }
    }
}
```

Run time was reduced to ~19.8s on Abacus.

Later Optimisation for Assignment2

- on Abacus with total core 4 : ~8.310297s (on average)
- on Abacus with total core 5 : ~6.212121s (on average)
- On my machine : ~5.6966s (on average)

Process for multicore programming optimisation

The OpenMP API uses the fork-join model of parallel execution. Multiple threads perform tasks defined implicitly or explicitly by OpenMP directives. All OpenMP applications begin as a single thread of execution, called the initial thread. The initial thread executes sequentially, until it encounters a parallel construct. At that point, this thread creates a group of itself and zero or more additional

threads and becomes the master thread of the new group. Each thread executes the commands included in the parallel region, and their execution may be differentiated, according to additional directives provided by the programmer. At the end of the parallel region, all threads are synchronized

```
void floyd_v1(int * matrix, int V)
{
    register int i,j,k;
    for(k=0;k<V;++k)
    {
        register int* km = (matrix + k*V);

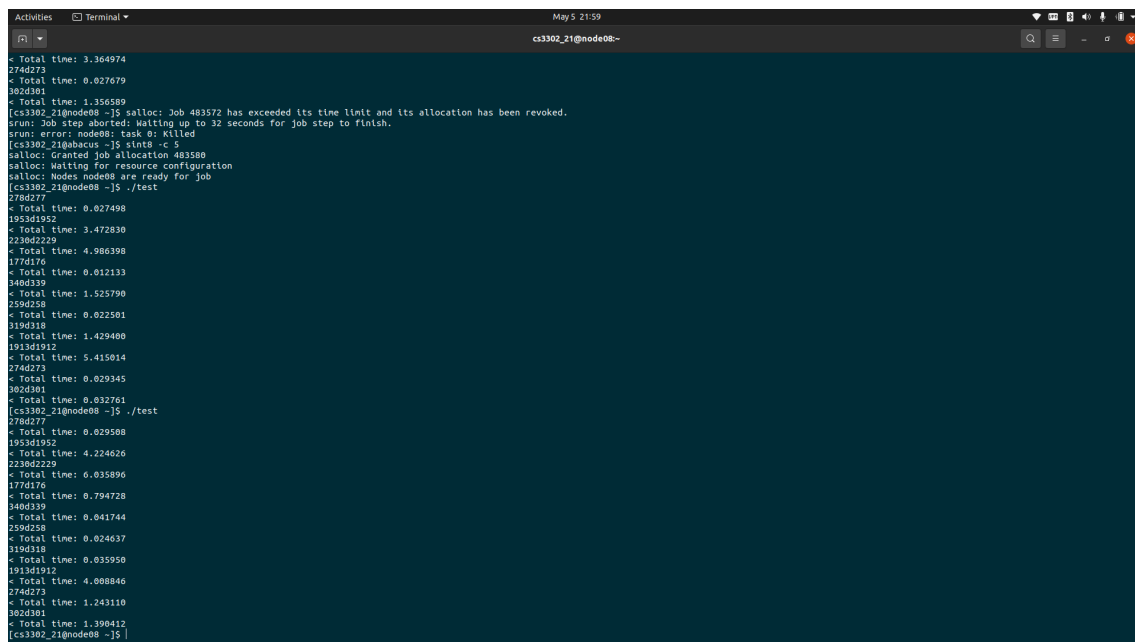
        #pragma omp parallel for private(j) shared(matrix,V)

        for(i=0;i<V;++i)
        {
            register int* im = (matrix + i*V);

            for(j=0;j<V;++j)
            {
                if( (*(im + j)) > ( (*(im + k) ) + ( *(km + j) ) ) )
                    (*(im + j)) = ( (*(im + k) ) + ( *(km + j) ) );
            }
        }
    }
}
```

Final Result

On Abacus we get best running time to be ~5s with our final code with core 5.



```
Activities Terminal May 5 21:59
cs3302_21@node08~
< Total time: 3.364974
274d273
< Total time: 0.027679
302d301
< Total time: 1.356589
[cs3302_21@node08 ~]$ salloc: Job 483572 has exceeded its time limit and its allocation has been revoked.
srun: Job step aborted: Waiting up to 32 seconds for job step to finish.
srun: error: node08: task 0: Killed
[cs3302_21@abacus ~]$ sinit8 -c 5
salloc: Granted job allocation 483588
salloc: Waiting for resource configuration
salloc: Nodes node08 are ready for job
[cs3302_21@node08 ~]$ ./test
278d277
< Total time: 0.027498
1953d1952
< Total time: 3.472830
2230d2229
< Total time: 4.986398
177d176
< Total time: 0.012133
340d339
< Total time: 1.525790
259d258
< Total time: 0.022501
119d118
< Total time: 1.429400
1913d1912
< Total time: 5.415014
274d273
< Total time: 0.029345
302d301
< Total time: 0.032761
[cs3302_21@node08 ~]$ ./test
278d277
< Total time: 0.029508
1953d1952
< Total time: 4.224626
2230d2229
< Total time: 6.035896
177d176
< Total time: 0.794728
340d339
< Total time: 0.041744
259d258
< Total time: 0.024637
119d118
< Total time: 0.035950
1913d1912
< Total time: 4.008846
274d273
< Total time: 1.243110
302d301
< Total time: 1.390412
[cs3302_21@node08 ~]$
```

Profiling on abacus:

gprof

```
[cs3302_21@node06 ~]$ cat gprof_work.sh
#!/bin/bash
gcc -pg -fopenmp -O2 code.c -o floyd
./floyd < testcases/Q2/t29 > temp
diff temp testcases/Q2out/t29
gprof -b ./floyd gmon.out
[cs3302_21@node06 ~]$ ./gprof_work.sh
2230d2229
< Total time: 5.149734
Flat profile:

Each sample counts as 0.01 seconds.
 %   cumulative   self           self       total
time  seconds    seconds   calls   ms/call  ms/call  name
100.97    26.17    26.17      1919     13.64    13.64  frame_dummy
  0.04     26.18     0.01                ms/call  ms/call  floyd_v1

Call graph

granularity: each sample hit covers 2 byte(s) for 0.04% of 26.18 seconds

index % time    self  children    called    name
-----
[1]   100.0     0.01   26.17      1919/1919  <spontaneous>
                                     floyd_v1 [1]
                                     frame_dummy [2]
-----
                                     floyd_v1 [1]
[2]   100.0    26.17     0.00      1919      frame_dummy [2]
-----

Index by function name

    [1] floyd_v1
    [2] frame_dummy
[cs3302_21@node06 ~]$ |
```

Valgrind

```
Total time: 502.608561
==232890==
==232890== I   refs:      100,931,167,470
==232890== I1  misses:      2,031
==232890== LLi misses:      2,007
==232890== I1  miss rate:      0.00%
==232890== LLi miss rate:      0.00%
==232890==
==232890== D   refs:      35,960,940,773 (35,301,967,179 rd + 658,973,594 wr)
==232890== D1  misses:      693,564,677 ( 693,131,478 rd + 433,199 wr)
==232890== LLd misses:      660,051,712 ( 659,628,559 rd + 423,153 wr)
==232890== D1  miss rate:      1.9% ( 2.0% + 0.1% )
==232890== LLd miss rate:      1.8% ( 1.9% + 0.1% )
==232890==
==232890== LL refs:      693,566,708 ( 693,133,509 rd + 433,199 wr)
==232890== LL misses:      660,053,719 ( 659,630,566 rd + 423,153 wr)
==232890== LL miss rate:      0.5% ( 0.5% + 0.1% )
```